

# DATA SCIENCE FESTIVAL

## 10 Data Science Uberhacks to Turbocharge your Workflow!!

Codenode, London  
20/05/2023



Tom Ewing  
Head of AI & Data Engineering  
Station10



tom-s10



@TomEwingS10


Station10

Firstly... sorry about the Clickbaity title.



Things this presentation was **nearly** called...

# Big Tech **HATES** these 10 Data Science Hacks!



Things this presentation was nearly called...

# 10 Reasons why your Data Science Workflow **SUCKS!**



Things this presentation was **nearly** called...

Multi-zillionaire tech-ninja open sources the  
**10 secrets** that made him his **FORTUNE!**



# About Me.

---



**Tom Ewing**

Head of AI & Data Engineering

Station10

- + 20 years experience in Data
- + 8 years experience in Data Science, Data Engineering & Machine Learning Engineering
- + Hands-on practitioner
- + I love “hacks”
- + Contributor to DSF since 2016



# What **Station10** do.



## Digital Analytics

Implementing tried & tested tools and strategies to maximise the earning potential of your website.



## Strategy & Insight

Revealing patterns and insights to better support data-driven decision making.



## AI & Data Engineering

Creating and embedding AI & data services and analysis to drive change, efficiency and value.


*“Creating value from data”*





# What this presentation is.

- + Gave a talk last year on Dask...
- + A lot of people there didn't know about PySpark's game-changing pandas-style API!





# What this presentation **is**.

- + This got me thinking... What else might people not know about?
- + Tech moves quickly!
- + There's things I like and use, but I can't fill 30 mins talking about one of them...  
So, maybe talk about all of them?!



# The rules.

I've come up with 10 “things” that:

- + Can make specific jobs or tasks around AI easier
- + You might not have heard of, or not be using
- + Are easy to pick up and integrate.
- + Are (mostly) Python-based.
- + Can showcased in 2 minutes or less.



# What this presentation **is NOT**.

- + A “deep dive” into anything.
- + About “hacking” outcomes or circumventing established AI processes.
- + Anything to do with Chat-GPT.




# Introducing the Uberhacks!



**1. Github Awesomeness**  
Awesome curated content.



**2. D-Tale**  
Eyeball data easily.



**3. Ydata-profiling**  
EDA as a Service.



**4. thefuzz**  
Fuzzy string matching.



**5. UK Open Data**  
So. Much. Data.



**6. Yellowbrick**  
Easy AI Visualisation.



**7. Shap**  
AI explainability.



**8. Fairlearn**  
Non-discriminatory AI.



**9. Metaflow**  
Easy Pipelines.



**10. Make**  
CLI maker-easier.





# **1. Github Awesomeness**

Awesome curated content.



# 1. Github **Awesomeness** Intro.

---



**Because Google is always trying to sell you stuff.**

- + Community curated content around a subject area (e.g. Machine Learning, Python, Data Engineering etc.) on Github.
- + Contain a list of links to packages, repos, sites, research papers... resources!
- + Wide ranging
- + Example: [awesome-production-machine-learning](#)





# 1. Github **Awesomeness** pros & cons.

---

## The Good...

- ✓ A great starting point for research
- ✓ Open Source centric
- ✓ Tend to be exhaustive (if updated)
- ✓ Quicker than Googling
- ✓ Opportunity to find more Uberhacks!

## The not so Good...

- ✗ Some topics are repeated across different lists.
- ✗ Some lists aren't updated regularly, so be sure to check.







## 2. D-Tale

Eyeball data easily.



## 2. D-Tale Intro.

---



**Excel but way better.**

- + Python package that generates an interactive dashboard to explore your data from your Notebook.
- + Has a range of exploration, visualisation and analysis options, tailored for ML use cases.
- + Viewed in the browser or a notebook.

Example: [Live D-Tale demo](#) (House Prices)



## 2. D-Tale example.



### Installation (CLI)

```
conda install dtale
```

### Execution (Python)

```
import dtale  
  
d = dtale.show(df)  
d.open_browser()
```



## 2. D-Tale pros & cons.

---



### The Good...

- ✓ Blazing fast as it loads what it needs dynamically
- ✓ View the whole dataframe not just the `.head()`
- ✓ Great breadth of functionality
- ✓ Does charting & dashboarding
- ✓ Very easy to use

### The not so Good...

- ✗ Doesn't work well in the cloud.
- ✗ Non-Python syntax for filtering which takes some getting used to.
- ✗ Crashes sometimes, particularly on larger datasets.
- ✗ Lots of dependencies






### 3. Ydata-profiling EDA as a Service.



### 3. Ydata-profiling Intro.

---




#### EDA done for you

- + The package formerly known as pandas-profiling
- + D-Tale = low level, Ydata-profiling = high level.
- + Shows data on duplicates, missing values, aggregations, correlations, statistics etc.

Example: [Live Ydata-profiling example \(Titanic\)](#)



### 3. Ydata-Profiling example.



#### Installation (CLI)

```
conda install ydata-profiling
```

#### Execution (Python)


```
From ydata_profiling import ProfileReport  
  
profile = ProfileReport(df)  
profile.to_file("my_report.html")
```





### 3. Ydata-profiling pros & cons.

---



#### The Good...

- ✓ Also, two lines of code!
- ✓ Saves time plotting and correlating
- ✓ Simple & explorable with a sharable output
- ✓ Has modes for dataset comparison and time series too.
- ✓ Makes it look like you've done more work than you actually have 😎

#### The not so Good...

- ✗ Slow particularly on larger datasets (but can now run on Spark)
- ✗ Limited scope





#### 4. thefuzz

Fuzzy string matching.



## 4. **thefuzz** Intro.

---



### Easy fuzzy joins in Pandas (or anything)

- + Python package that measures and matches similarity between two string elements.
- + Wide range of matching options including comparison, partial and token (sentence) matching
- + Can be used to join two dataframes with similar but non-identical keys



## 4. thefuzz examples.



### Installation (CLI)

```
conda install thefuzz
```

### Execution (Python)

```
from thefuzz import fuzz
from thefuzz import process

>>> fuzz.ratio("this is a test", "this is a test!")
97

>>> fuzz.partial_ratio("this is a test", "this is a test!")
100

>>> fuzz.token_sort_ratio("fuzzy wuzzy was a bear", "wuzzy fuzzy was a bear")
100

>>> choices = ["Atlanta Falcons", "New York Jets", "New York Giants", "Dallas Cowboys"]
>>> process.extract("new york jets", choices, limit=2)
[('New York Jets', 100), ('New York Giants', 78)]
```



## 4. thefuzz pros & cons.

---



### The Good...

- ✓ Simple & effective
- ✓ Quick (considering...)
- ✓ Makes “dirty joins” possible with a degree of control
- ✓ Versatile and can be used outside of pandas

### The not so Good...

- ✗ Slow if your data is large
- ✗ Need to be careful and check the joins





## 5. UK Open Data So. Much. Data.



## 5. UK Open Data Intro.

---

### Open data can make your stuff better!

- + UK has a wide range of open data particularly from central statistical authorities and government (ONS, Gov.Scot, NIRSA, Gov.uk etc.)
- + A multitude of areas (economic, demographic, housing, deprivation, transport etc.) that can enhance outcomes.
- + 2021 / 2022 Census data is presently “fresh”





## 5. ONS Open Data pros & cons.

### The Good...

- ✓ Enables per-capita comparisons between areas
- ✓ Huge breadth of data
- ✓ Deep granularity
- ✓ A range of geographic breakdowns
- ✓ Well supported by the ONS Geoportal

### The not so Good...

- ✗ ONS website isn't user friendly
- ✗ Need some geographic knowledge to merge to existing data
- ✗ Time consuming to collate and maintain
- ✗ Disparate sources across England & Wales, Scotland and Northern Ireland





## 6. Yellowbrick

Easy AI Visualisation.



## 6. Yellowbrick Intro.

---












**Because C+P matplotlib is so 2017...**

- + Python package that visualises AI algorithms.
- + Extends Scikit-Learn
- + Covers regression, classification, clustering and much more.



## 6. Yellowbrick gallery.



## 6. Yellowbrick examples.



### Installation (CLI)

```
conda install -c districtdatalabs yellowbrick
```

### Execution (Python)

```
from yellowbrick.classifier import ClassificationReport
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
visualizer = ClassificationReport(model)

visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
visualizer.show()
```



## 6. Yellowbrick pros & cons.



### The Good...

- ✓ Beautiful
- ✓ Wide range of visualisers
- ✓ Simple & works out of the box
- ✓ Excellent prototyping tool
- ✓ Excellent documentation & examples gallery

### The not so Good...

- ✗ SKL only
- ✗ “Wrapper” may cause issues if using in conjunction with other wrappers
- ✗ Sometimes you have to train twice
- ✗ Visualisers aren’t easily customisable





## 7. Shap

AI Explainability.





## 7. **Shap** Intro.

---



### **Explainability isn't just for stakeholders...**


- + Uses game theory to explain the predictions of any model.
- + Calculates the contribution of each feature for each record in the overall prediction data.
- + This uncovers complex relationships between the features and the target and uncovers how models are working “under the hood”.
- + Reveals where features do well and badly in the data.



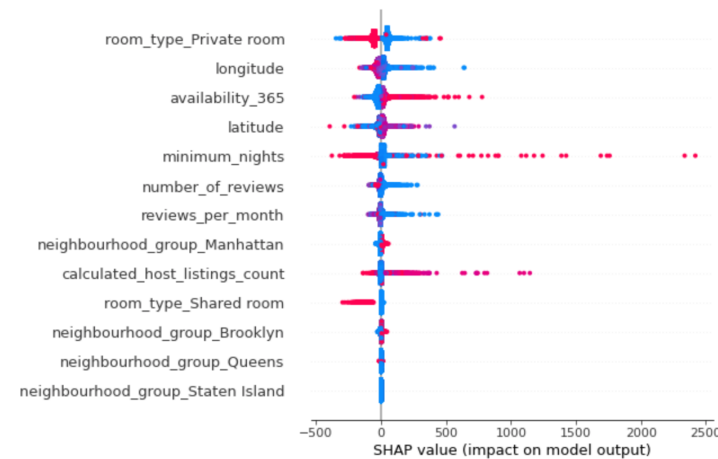
# 7. Shap gallery.




## Waterfall Plot



## Summary Plot



## Force Plot



Red = Positive contribution  
Blue = Negative contribution



## 7. Shap examples.



### Installation (CLI)

```
conda install shap
```

### Execution (Python)

```
# Initialise the explainer & Calculate Shap values
explainer = shap.Explainer(model, X_train)
shap_values = explainer(X_test)

# Waterfall plot on a single record
shap.plots.waterfall(shap_values[0])

# Summary plot for all records
shap.summary_plot(shap_values, X_test)

# Force plot
shap.force_plot(explainer.expected_value, shap_values, X_test)
```



## 7. **Shap** pros & cons.

---



### The Good...

- ✓ Simple to use...
- ✓ Works with most major AI libraries
- ✓ Interactive plotting in your notebook
- ✓ Beautiful
- ✓ Excellent for images

### The not so Good...

- ✗ Complex to master!
- ✗ Highly mathematical
- ✗ Computationally expensive





## 8. Fairlearn

Non-discriminatory AI.



## 8. Fairlearn Intro.

---



### AI Regulation is coming!

- + Python package that helps identify and mitigate unfairness in AI models.
- + Has two components:
  - + **Metrics** for identifying which groups are negatively impacted by a model
  - + **Algorithms** for mitigating any unfairness



## 8. Fairlearn examples.



### Installation (CLI)

```
conda install fairlearn
```

### Execution (Python)

```
# Metric (Gender)

from fairlearn.metrics import MetricFrame
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier

classifier = DecisionTreeClassifier(min_samples_leaf=10, max_depth=4).fit(X, y_true)

y_pred = classifier.predict(X)
mf = MetricFrame(
    metrics=accuracy_score,
    y_true=y_true,
    y_pred=y_pred,
    sensitive_features=sex
)

mf.overall
mf.by_group
```



## 8. Fairlearn examples.



### Execution (Python)

```
# Algorithm (Gender)

from fairlearn.reductions import (
    ExponentiatedGradient,
    DemographicParity
)

constraint = DemographicParity()
classifier = DecisionTreeClassifier(min_samples_leaf=10, max_depth=4)

mitigator = ExponentiatedGradient(classifier, constraint).fit(X, y_true, sensitive_features=sex)

y_pred_mitigated = mitigator.predict(X)

sr_mitigated = MetricFrame(metrics=selection_rate, y_true=y_true, y_pred=y_pred_mitigated, sensitive_features=sex)

sr_mitigated.overall
sr_mitigated.by_group
```





## 8. Fairlearn pros & cons.

---



### The Good...

- ✓ Integrates with major AI packages (SKL, PyTorch etc.)
- ✓ Wide range of fairness metrics & algorithms
- ✓ Great documentation
- ✓ Integrated into Azure
- ✓ It's ethical and the right thing to do

### The not so Good...

- ✗ Another “Wrapper”
- ✗ Highly mathematical
- ✗ Computationally expensive
- ✗ Algorithms need data on gender, ethnicity, religion, immigration etc.
- ✗ Be prepared to take a hit on model performance





## 9. Metaflow

Easy Pipelines.



## 9. Metaflow Intro.

---



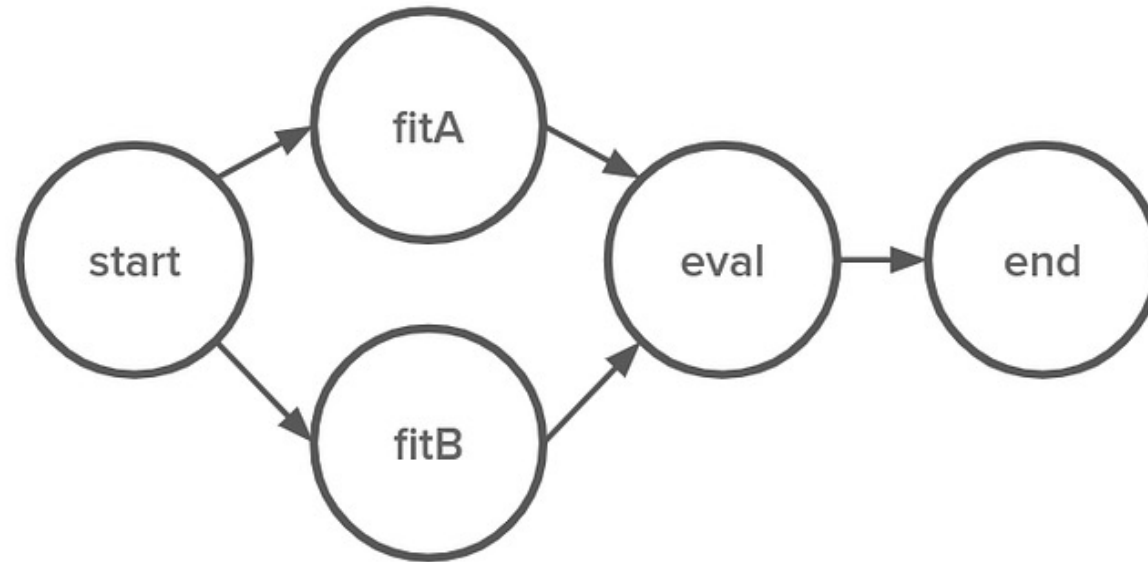
**Because everyone loves pipelines!**

- + Netflix produced “Data Science Project Manager”
- + Versatile and easy orchestration tool:
  - + In-project pipelines
  - + End-to-end data workflows
  - + Environment management



## 9. Metaflow examples.

**The DAG (Directed Acyclic Graph):**



## 9. Metaflow examples.



### Installation (CLI)

```
conda install metaflow
```

### Execution (Python)

```
from metaflow import FlowSpec, step

class MyFlow(FlowSpec):

    @step
    def start(self):
        self.data = load_your_data()
        self.next(self.step_a, self.step_b)

    @step
    def fitA(self):
        self.model= model_a()
        self.next(self.eval)

    @step
    def fitB(self):
        self.model=model_b()
        self.next(self.eval)

    @step
    def eval(self, inputs):
        self.best = pick_best(inputs)
        self.next(self.end)

    @step
    def end(self):
        logger.success("All done!")
```



## 9. Metaflow pros & cons.

---



### The Good...

- ✓ Simple, versatile & scalable
- ✓ Easy parallelisation of workflows
- ✓ Extensible “microframework”
- ✓ GUI option

### The not so Good...

- ✗ Verbose
- ✗ Limited Cloud support for non-AWS





**10. Make**  
CLI maker-easier.



## 10. Make Intro.



### CLI maker-easier

- + Automates the process of building programs through managing dependencies (Think data pipelines but for programs!)
- + It also enables creation of repeatable CLI jobs (what I use it for)
- + Specify complex tasks in the **Makefile** and run these with `make <command>`





# 10. Make examples.



## Make installation

**Mac:** Included in Xcode

**Windows:** Via chocolate package manager

**Linux:** Included already!

## Makefile example

```
CONDA_ENVIRONMENT_NAME = my_env
PYTHON_VERSION = 3.8

ACTIVATE = source $$($(conda info --base)/etc/profile.d/conda.sh ; conda activate ; conda activate
DEACTIVATE = source $$($(conda info --base)/etc/profile.d/conda.sh ; conda deactivate ; conda deactivate

# Environment Management
.PHONY: create-environment
create-environment: ## Create the conda environment & jupyter kernel
conda create --name $(CONDA_ENVIRONMENT_NAME) --channel conda-forge --yes python=$(PYTHON_VERSION)
$(ACTIVATE) $(CONDA_ENVIRONMENT_NAME) && \
    conda config --add channels conda-forge && \
    conda config --set channel_priority strict && \
    conda install --yes --file requirements-conda.txt && \
    conda update pip && \
    pip install -r requirements-pip.txt && \
    conda env export --name $(CONDA_ENVIRONMENT_NAME) --no-builds > environment.yaml && \
    ipython kernel install --user --name $(CONDA_ENVIRONMENT_NAME) && \
    nbstripout --install --attributes .gitattributes
$(DEACTIVATE)

.PHONY: remove-environment
remove-environment: ## Remove the environment and any relevant files
conda env remove --name $(CONDA_ENVIRONMENT_NAME)

.PHONY: run
run: ## Run the application
$(ACTIVATE) $(CONDA_ENVIRONMENT_NAME) && \
    python -m main --env=.env
$(DEACTIVATE)
```



## 10. Make pros & cons.



### The Good...

- ✓ Simplifies complex CLI jobs
- ✓ Pre-installed on most VMs
- ✓ Easy model serving via REST API
- ✓ “Hackable” to an extent
- ✓ Made in 1976... Stood the test of time

### The not so Good...

- ✗ Verbose & complex syntax
- ✗ Unforgiving (e.g. requires tabs not spaces)
- ✗ Difficult to install on windows machines



# 11. Bonus Round!

---

The honourable mentions:

- + [polars](#): High Performance Computing (HPC) in Python
- + [dask](#): Like Spark but on your machine
- + [splink](#): Data linking at scale
- + [ydata-synthetic](#): Synthetic data generator
- + [python-dp](#): Python data privacy
- + [scikit-plot](#): AI explanation & visualisation
- + [lime](#): Quicker explainable AI
- + [pyLDAvis](#): Interactive topic model visualisation
- + [mlflow](#): AI tracking & serving
- + [nbstripout](#): Remove Jupyter output cells from git
- + [loguru](#): Easy, colourful logging
- + [Excalidraw](#): Online collaborative whiteboarding
- + [Github copilot \(\\$\)](#): Predictive coding AI
- + [DuckDB](#): SQLite for analytics
- + [Streamlit](#): R Shiny! In Python
- + [SparseSC](#): Synthetic A/B testing




## 12. Bonus Round #2.

---

The not-so honourable mentions from Reddit:

- + Domain knowledge
- + Communication skills
- + Sense of humour
- + Economics
- + Psychometrics
- + My data engineer
- + Pandas
- + Excel
- + Powerpoint
- + The side scroller on my mouse



# Those 10 **Uberhacks** again.




## 1. Github Awesomeness

<https://github.com>



## 2. D-Tale

<https://github.com/man-group/dtale>



## 3. Ydata-profiling

<https://github.com/ydataai/ydata-profiling>



## 4. Thefuzz

<https://github.com/seatgeek/thefuzz>



## 5. UK Open Data

<https://www.ons.gov.uk/>  
<https://opendata.scot/>  
<https://www.nisra.gov.uk/>



## 6. Yellowbrick

<https://github.com/DistrictDataLabs/yellowbrick>



## 7. Shap

<https://github.com/slundberg/shap>



## 8. Fairlearn

<https://github.com/fairlearn/fairlearn>



## 9. Metaflow

<https://metaflow.org>



## 10. Make

<https://www.gnu.org/software/make>



# Thanks!

Slides, links, posts & mo' shizzle on my socials from Monday!

 tom-s10

 @TomEwingS10

Station10

[www.station10.co.uk](http://www.station10.co.uk)



?

Station10

[www.station10.co.uk](http://www.station10.co.uk)